

COMP 252: Summary

Anna Brandenberger

April 28, 2018

This is a condensed scribing of the lectures given by Luc Devroye in the Winter semester of 2018 for the Honours Data Structures and Algorithms class (COMP 252, McGill University). Subjects covered: Divide-and-Conquer algorithm examples, Dynamic Programming, various Abstract Data Types and implementations.

Definitions & Theorems

Introduction

Definition 1. An **algorithm** is a finite set of instructions which takes in a finite number of inputs, performs a task exactly, produces a finite output and exits. (c.f. Church-Turing Thesis.)

Definition 2. We examine two **cost models**:

1. *Uniform cost model*: every operation takes 1 time unit.
2. *Bitwise (logarithmic) cost model*: every *bit* operation takes 1 time unit (i.e. basic operations like reading, shifting, etc.).

Complexity

Definition 3. An **oracle** is a black-box type device capable of solving some computational problem in a single time unit: takes in input(s) x_1, \dots, x_n and provides output(s) y_1, \dots, y_k .

Definition 4. **O - Big O**: upper bound. We say that $a_n = O(b_n)$ if $\exists c > 0, n_0 > 0$ such that $a_n \leq cb_n, \forall n \geq n_0$.

Definition 5. **Ω - Big Omega**: lower bound. We say that $a_n = \Omega(b_n)$ if $\exists c > 0, n_0 > 0$ such that $a_n \geq cb_n, \forall n \geq n_0$.

Definition 6. **Θ - Big Theta**: precise asymptotic behaviour. We say that $a_n = \Theta(b_n)$ if $a_n = O(b_n) = \Omega(b_n)$.

Lower Bounds chapter: if the lower bound for solving a problem with input of size n is f_n , it implies that \forall algorithms (most efficient), \exists an input x_1, \dots, x_n (worst-case) s.t. $T(\text{algorithm}, x_1, \dots, x_n) \geq f_n$.

Recurrences

Theorem 7 (Master Theorem). Given a recurrence $T_n = aT_{\lfloor n/b \rfloor} + f(n)$,

(a) $n^{\log_b a} / f(n) > n^\epsilon \quad \forall n \geq n_0 \implies T_n = O(n^{\log_b a})$.¹

(b) $f(n) / n^{\log_b a} > n^\epsilon \quad \forall n \geq n_0 \implies T_n = O(f(n))$.²

(c) $f(n) = \Theta(n^{\log_b a}) \implies T_n = \Theta(n^{\log_b a} \cdot \log_b n)$.

¹ If $n^{\log_b a} / f(n) = \Theta(n^c) \quad \forall n \geq n_0$, then $T_n = \Theta(n^{\log_b a})$ instead.

² This case holds under the technical assumption

$$\lim_{n \rightarrow \infty} \left(\frac{f(n)}{a \cdot f(n/b)} \right) > 1.$$

Definition 8. An algorithm's **recursion tree** is an a -ary tree used to visualize its recurrence given as $T_n = a \cdot T_{n/b} + f(n)$. Calling T_n costs $f(n)$ units of "work" and results in a new problem with worst-case time $T_{n/b}$ being called " a " times. Each node in the tree represents an algorithm call, with the root node referring to the first call on a problem of size n . Given the recursion tree for some algorithm, the **time complexity** of the algorithm is the sum of the required work over all of its nodes.

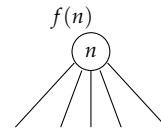
Lower Bounds

c.f. Devroye's printed notes for theory.

Data Structures

Definition 9. An **endogenous** data structure stores values in the structure, while an **exogenous** data structure maintains pointers to externally stored values.

The node representing a problem of time n is denoted by



where $f(n)$ is the work required, and each of the " a " lines lead to a problem of size n/b .

1 Divide-and-Conquer

LECTURE 01/09

Example 1.1. Fast Exponentiation: compute x^n efficiently.

$$\begin{bmatrix} x_n \\ x_{n-1} \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} x_{n-1} \\ x_{n-2} \end{bmatrix} = M^2 \begin{bmatrix} x_{n-2} \\ x_{n-3} \end{bmatrix} = \dots = M^{n-1} \begin{bmatrix} x_1 \\ x_0 \end{bmatrix}$$

TIME COMPLEXITY:

RAM model:

$$T_n = T_{\frac{n}{2}} + 1$$

$$\Rightarrow T_n = \Theta(\log_2 n)$$

Bit model:

$$T_n = T_{\frac{n}{2}} + n^2$$

$$\Rightarrow T_n = \Theta(n^2)$$

LECTURE 01/11 ³

Example 1.2. Chip Testing: determine state (G/B) of all chips, where tester oracle has two chips evaluate each other, i.e. 3 possible outputs: GG (both from \mathcal{G} or both from \mathcal{B}), GB and BB (at least one from \mathcal{B}).⁴

ALGORITHM(n)

- 1 GOOD = FINDGOODCHIP(n)
- 2 do $n - 1$ tests against GOOD to determine state of all chips.

FINDGOODCHIP(n)

- 1 if $n = 1$ return the one chip.
- 2 if n even
 - 3 pair all chips and test ($n/2$ pairs).
 - 4 eliminate all non-GG pairs.
 - 5 eliminate the second chip of each GG pair.
 - 6 FINDGOODCHIP(remaining chips)
- 7 else if n odd
 - 8 take a random chip x and test against all other chips.
 - 9 take a majority vote:
 - 10 if $\geq \frac{n-1}{2}$ G votes: $x \in \mathcal{G}$
 - 11 return x
 - 12 else $x \in \mathcal{B}$
 - 13 eliminate x
 - 14 FINDGOODCHIP(remaining chips with n even)

TIME COMPLEXITY of FINDGOODCHIP(n):

$$\begin{cases} T_n \leq \frac{3n}{2} + T_{\lfloor \frac{n}{2} \rfloor} \\ T_1 = 0, T_2 = 0 \end{cases} \Rightarrow T_n = \Theta(n)$$

Compute M^n , where $M = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}$:MATRIX(n)

- 1 if $n = 0$ return $\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
- 2 if $n = 1$ return M
- 3 else
 - 4 if n even return MATRIX($\frac{n}{2}$)
 - 5 if n odd return MATRIX($\frac{n-1}{2}$) $\cdot M$

³ Scribed: 2018.

⁴ Given $|\mathcal{G}| > |\mathcal{B}|$ where:
 \mathcal{G} : set of good chips;
 \mathcal{B} : set of bad chips

Example 1.3. Karatsuba Multiplication: multiply two n -bit numbers a_n and b_n efficiently.

$$a_n \times b_n = \alpha_1 \beta_1 + \alpha_2 \beta_2 + (\alpha_1 \beta_2 + \alpha_2 \beta_1) \times 2^{n/2}$$

where $\alpha_1 \beta_2 + \alpha_2 \beta_1 = (\alpha_1 - \alpha_2)(\beta_1 - \beta_2) + \alpha_1 \beta_1 + \alpha_2 \beta_2$.

MULTIPLY(a_n, b_n)

- 1 ONE = MULTIPLY(α_1, β_1)
- 2 TWO = MULTIPLY(α_2, β_2)
- 3 THREE = MULTIPLY($\alpha_1 - \alpha_2, \beta_1 - \beta_2$)
- 4 **return** ONE + TWO + (ONE + TWO + THREE) $\times 2^{n/2}$ (shifting)

TIME COMPLEXITY:

$$T_n = 3T_{\frac{n}{2}} + 2n + 10n \Rightarrow T_n = \Theta(n^{\log_2 3})$$

Example 1.4. Convex Hull: find C.H. in order given n points in \mathbb{R}^2 .

FINDUPPERHULL($1 \dots n$)

- 1 FINDUPPERHULL($1, \dots, \frac{n}{2}$).
- 2 FINDUPPERHULL($\frac{n}{2} + 1, \dots, n$).
- 3 MERGE:
 - progress by x -coordinate from left to right,
 - eliminate points with angle $< 180^\circ$.

TIME COMPLEXITY:

$$T_n = 2T_{\frac{n}{2}} + n \Rightarrow T_n = \Theta(n \log n)$$

LECTURE 01/16 ⁵

Example 1.5. Strassen Matrix Multiplication: matrix multiply two $n \times n$ matrices A and B .

$$A \times B = \left[\begin{array}{c|c} A_1 & A_2 \\ \hline A_3 & A_4 \end{array} \right] \times \left[\begin{array}{c|c} B_1 & B_2 \\ \hline B_3 & B_4 \end{array} \right] = \left[\begin{array}{c|c} A_1 B_1 + A_2 B_2 & A_1 B_2 + A_2 B_4 \\ \hline A_3 B_1 + A_4 B_3 & A_3 B_2 + A_4 B_4 \end{array} \right]$$

Use algebraic trick to reduce the 8 multiplications to 7, yielding:

TIME COMPLEXITY:

$$T_n = n^2 + 7T_{n/2} \Rightarrow T_n = \Theta(n^{\log_2 7}) \approx \Theta(n^{2.807})$$

LECTURE 01/18 ⁶

Example 1.6. Selection Problem: finding the k th smallest number in a set $S = \{x_1, \dots, x_n\}$.

$$\begin{cases} a_n = (\underbrace{[010\dots]}_{\alpha_2} \underbrace{[\dots 001]}_{\alpha_1}) = \alpha_1 + \alpha_2 \times 2^{n/2} \\ b_n = (\underbrace{[000\dots]}_{\beta_2} \underbrace{[\dots 110]}_{\beta_1}) = \beta_1 + \beta_2 \times 2^{n/2} \end{cases}$$

TOOM-COOK OPTIMIZATION:

$$T_n = 5T_{n/3} + o(n) = \Theta(n^{\log_3 5})$$




Figure 1: Diagram of the merging step. The blue and black lines are the two upper hulls.

⁵ Scribed: 2018, Adam.

⁶ Scribed: 2018, Olivia

LECTURE 01/23

Example 1.7. Mergesort. Prove inductively that $T_n = \Theta(n \log n)$.

$T_n = n - 1 + T_{\lfloor \frac{n}{2} \rfloor} + T_{\lceil \frac{n}{2} \rceil}$, where n : number of comparisons.

Proof. Hypothesis $P(n)$: $\exists c > 0$ s.t. $T_n \leq cn \log_2 n$.

- **Base Case** $n = 0$: OK.
- **Inductive Step**: Assume $P(k)$ true up to $n - 1$.

<p>n even:</p> $ \begin{aligned} T_n &\leq n - 1 + 2c \left(\frac{n}{2}\right) \log_2 \left(\frac{n}{2}\right) \\ &= n - 1 + cn \log_2 n - cn \\ &\leq cn \log_2 n \quad \text{for all } c \geq 1 \end{aligned} $	<p>n odd:</p> $ \begin{aligned} T_n &\leq n - 1 + c \left(\frac{n+1}{2}\right) \log_2 \left(\frac{n+1}{2}\right) + c \left(\frac{n-1}{2}\right) \log_2 \left(\frac{n-1}{2}\right) \\ &= n - 1 + c \frac{n}{2} \log_2 \left(\frac{n^2 - 1}{4}\right) + \frac{1}{2} c \log_2 \left(\frac{n+1}{n-1}\right) \\ &\leq n - 1 + c \frac{n}{2} \log_2 \left(\frac{n^2}{4}\right) + \frac{1}{2} c \log_2(2) \\ &= n - 1 + c \frac{n}{2} (2 \log_2 n - 2) + \frac{1}{2} c \\ &= n - 1 - c \left(n - \frac{1}{2}\right) + cn \log_2 n \\ &\leq cn \log_2 n \quad \text{for all } c \geq 0 \quad \square \end{aligned} $
--	--

Example 1.8. Selection Problem Take II. Finding the k^{th} smallest element in a set of size n .

2 Dynamic Programming

Example 2.1. Binomial Coefficient.

Example 2.2. Travelling Salesman Problem.

LECTURE 01/25 ⁷⁷ Scribed: 2018.

Example 2.3. LCS.

Example 2.4. Knapsack Problem.

Example 2.5. Optimal Binary Search Tree

- **Context**: Static trees are simpler versions of normal dynamic search trees: no adding, removing, etc. They are used for example in CPU for parsing programs (reading key words).
The cost of the tree is defined as: $C = \sum_{i=1}^n w_i l_i$ where w_i are weights of each key and l_i its level (distance from the root)
- **Goal**: Given a static set of n ordered keys with weights assigned to each, find the minimizing tree (minimize C).

ALGORITHM to find $C[1, n]$:

$C[i, j]$ = minimal optimal cost for problem with items $i \dots j$.

\Rightarrow Will need $W[i, j] = \sum_{k=1}^j w_k = \text{weights of } i \dots j$.⁸

⁸ This can be done in $\Theta(n^2)$ time:

DYNAMIC PROGRAM

```

// Initialization
1 for i = 1 to n
2   C[i, i] = w_i      // tree with only one key
3   if j < i           // invalid tree (interval with end index > start index)
4     C[i, j] = 0
// Recursive Step
5 for size = 2 to n    // size of interval
6   for i = 1 to n      // starting index of interval
7     j = i + size - 1
8     if j ≤ n          // to ensure a valid interval (last index max n)
9       C[i, j] = min_{i ≤ k ≤ j} (C[i, k - 1] + C[k + 1, j] + W[i, j])
10      root[i, j] = one of the r's that minimizes C[i, j]
11 return C, root

```

WEIGHTS(w_i)

```

1 for i = 1 to n
2   W[i, i] = w_i
3   for j = i + 1 to n
4     W[i, j] = W[i, j - 1] + w_j

```

$$\begin{aligned}
 C[i, j] &= \sum_{\substack{i=1 \\ i \neq k}}^j w_i (l_i + 1) + w_{root} \\
 &= \sum_{\substack{i=1 \\ i \neq k}}^j w_i l_i + \sum_{\substack{i=1 \\ i \neq k}}^j w_i + w_{root} \\
 &= \sum_{i=1}^{k-1} w_i l_i + \sum_{s=k+1}^j w_s l_s + \sum_{t=i}^j w_t \\
 &= C[i, k - 1] + C[k + 1, j] + W[i, j]
 \end{aligned}$$

\Rightarrow The complexity of this algorithm is $O(n^3)$ since we have
 $\sum_{k=2}^n \sum_{m=1}^n \text{size computations/runs through the loop.}^9$

⁹ Knuth reduced this to n^2 : look up and add info.

LECTURE 01/30

Example 2.6. Matching Problem.

Example 2.7. Job Scheduling.

LECTURE 02/01 - LOWER BOUNDS ¹⁰

¹⁰ c.f. Devroye's printed notes.

3 Data Structures

LECTURE 02/06

3.1 Lists

LECTURE 02/08

3.2 Trees

types of trees
 traversals & transforming btw storing representations
 lung diagram

LECTURE 02/13

dictionaries

3.3 Red-Black Trees

LECTURE 02/20 ¹¹¹¹ Scribed: 2018, Akshal

3.4 Augmenting Data Structures

LECTURE 02/22 ¹²¹² c.f. CLRS

Example 3.1. RB tree + Linked List.

Example 3.2. Order Statistic Tree.

Example 3.3. Interval Tree.

Example 3.4. Level Linked RB Tree (for fast browsing).

LECTURE 02/27 ¹³¹³ Scribed: 2017 notes.

3.5 Cartesian Trees

Definition 10. A **cartesian tree** on n nodes (Figure 3.5) is a binary tree storing n tuples $(x_i, y_i) : (key, timestamp)$: we build a BST by inserting the nodes by increasing timestamps. The cartesian tree is therefore a *binary search tree* with respect to keys x_i , and a *heap* with respect to timestamps y_i .

Example 3.5. BST: inserting x_1, \dots, x_n in sequence, i.e. a cartesian tree for $(x_1, 1), (x_2, 2), \dots, (x_n, n)$.

Example 3.6. RBST: randomly permuting all elements before adding them, i.e. a cartesian tree for $(x_1, \tau_1), (x_2, \tau_2), \dots, (x_n, \tau_n)$ where (τ_1, \dots, τ_n) is a random permutation of $(1, \dots, n)$.

Example 3.7. Treap: same concept as RBST, i.e. a cartesian tree for $(x_1, u_1), (x_2, u_2), \dots, (x_n, u_n)$ where $u_i \in [0, 1] \in \mathbb{R}$ random.

ATOMIC OPERATIONS ON CARTESIAN TREES:

i.e. FIX-ing operations: the input is (v, t) , a Cartesian tree t in which only one node, v , does not satisfy the heap property.

FIX-UP(*node*)

```

1  while node ≠ root && node.y < node.parent.y
2      if node == node.parent.right then
3          node.parent.right = node.left
4          node.left = node.parent
5          node.parent.left = node.right
6          node.right = node.parent

```

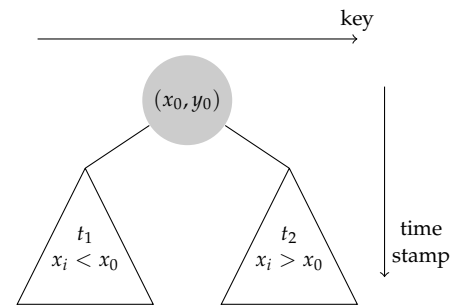


Figure 2: Diagram of a cartesian tree: (BST) keys are increasing when sweeping from left to right, and (heap) timestamps are increasing on any path from a root to a leaf.

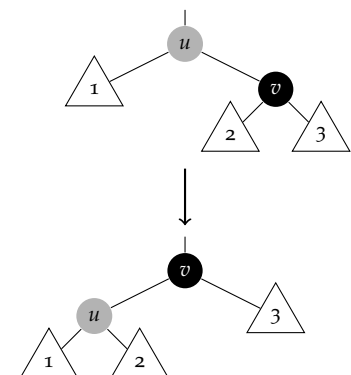


Figure 3: tree rotation done during FIX-UP and FIX-DOWN. In FIX-UP, the black node v is misplaced, it should be above the gray node u . In FIX-DOWN, the gray node is misplaced, it should be below the black one (only difference is which we are pointing to).

FIX-DOWN(*node*)

```

1  while node is not leaf && (node.y > node.left.y || node.y > node.right.y)
2      if node.right.y > node.left.y then
3          node.leftt = node.left.right
4          node.left.right = node
5          node.right = node.right.left
6          node.right.left = node

```

OTHER OPERATIONS ON CARTESIAN TREES:

Exercise 3.8. Study how to do these (recursively) without using FIX-UP and FIX-DOWN.

1. INSERT($t, (x, y)$): Inputs are a Cartesian tree t and a node (x, y) to be added to the tree. First, insert the node (x, ∞) just as you would do it in a binary search tree, implying the node will be a leaf since the heap property must be satisfied. Second, change the node to (x, y) and use FIX-UP to fix the Cartesian tree.
2. DELETE($t, node$): Inputs are a Cartesian tree t and a pointer $node$ to a node that will be removed. First, change $node.y$ to ∞ and use FIX-DOWN to move the node down the tree. It will end up as a leaf so the second step is just to remove the node.
3. JOIN(t_1, t_2): Inputs are two Cartesian trees t_1 and t_2 that need to be joined into a single tree, given that all keys in t_1 are smaller than all keys in t_2 . First, create a temporary node $(k, -\infty)$ such that $MAX_X(t_1) < k < MIN_X(t_2)$ and let its left child be the root of t_1 while its right child is the root of t_2 . Second, delete the node and the tree will fix itself up.
4. SPLIT(t, k): Inputs are a Cartesian tree t and a value k that will split the tree into two trees that have x coordinates smaller than k and bigger than k , respectively. First, insert a temporary node $(k, -\infty)$ (it will end up as a root) and after the procedure, the left subtree and right subtree of the node are the trees we are looking for.

EXPECTED VALUE ANALYSIS FOR TREAPS:

Define D_k to be the depth of the node with rank k , namely (k, T_k) . Since the depth of a node is equivalent to the number of ancestors it has, let

$$D_k = \sum_{j \neq k} X_{jk} \text{ where } X_{jk} = \begin{cases} 1, & \text{if } j \text{ ancestor of } k \\ 0, & \text{otherwise} \end{cases}$$

So next we wish to calculate the expected value for D_k .

$$\begin{aligned}
 \mathbb{E}[D_k] &= \sum_{j \neq k} \mathbb{P}((j, T_j) \text{ is an ancestor of } (i, T_i)) \\
 &= \sum_{j \neq k} \mathbb{P}(T_j \text{ is the smallest of } T_j, T_{j+1}, \dots, T_k) \\
 &= \sum_{j < k} \frac{1}{k - j + 1} + \sum_{j > k} \frac{1}{j - k + 1} \\
 &= \left(\frac{1}{2} + \dots + \frac{1}{k} \right) + \left(\frac{1}{2} + \dots + \frac{1}{n - k} \right) \\
 &= (H_k - 1) + (H_{n-k+1} - 1) \\
 &\leq 2 \ln(n) \simeq 1.39 \log_2(n).
 \end{aligned}$$

We will denote H_k to be the harmonic number, $\sum_{n=1}^k \frac{1}{n} = 1 + \frac{1}{2} + \dots + \frac{1}{k}$. It can be approximated by $\int_1^k \frac{1}{x} dx \approx \ln(k)$ but more importantly, it can be bounded as $H_k \in [\ln(k+1), \ln(k+1) + 1]$.

Example 3.9. Quicksort Tree: the QUICKSORT procedure is equivalent to building a Treap (or RBST): same number of comparisons needed.

QUICKSORT(*list*, *low*, *high*)

- 1 i = random index between *low* and *high*
- 2 $a = \text{list}[i]$
- 3 **if** *high* > *low* **then**
- 4 Partition *list*[*low*...*high*] so elements before index j are smaller than a and elements after index j are greater than a
- 5 QUICKSORT(*list*, *low*, $j - 1$)
- 6 QUICKSORT(*list*, $j + 1$, *high*)

EXPECTED COMPLEXITIES FOR QUICKSORT:

1. Let C_n = # of comparisons = $\sum_{i=1}^n D_i$.

$$\begin{aligned}
 \mathbb{E}[C_n] &= \sum_{i=1}^n \mathbb{E}[D_i] = \sum_{i=1}^n (H_i + H_{n-i+1} - 2) = 2 \sum_{i=1}^n (H_i - 1) \\
 &= 2 \sum_{i=1}^n H_i - 2n = 2 \sum_{i=1}^n \sum_{j=1}^i \frac{1}{j} - 2n = 2 \sum_{j=1}^n \frac{1}{j} \sum_{i=j}^n 1 - 2n \\
 &= 2 \sum_{j=1}^n \frac{1}{j} (n - j + 1) - 2n = 2(n+1)H_n - 4n \\
 &\simeq 2n \ln(n) = 1.386294 \dots n \log_2(n)
 \end{aligned}$$

2. Operations INSERT, DELETE, JOIN and SPLIT are all $O(\log n)$.

LECTURE 03/01

3.6 Priority Queues

Definition 11. A **priority queue** is an abstract data structure storing items with keys, supporting operations INSERT, DELETEMIN,

DELETE, DECREASEKEY and sometimes SEARCH. Implementation possibilities are red-black trees, standard heaps or **tournament trees**. It can be applied to (1) sorting, (2) operating systems, and (3) Discrete Event Simulation (DES).

Definition 12. A **tournament tree** is a complete binary tree where all items are leaves (n) and all internal nodes represent matches between two items, where they store pointers to winners of the match ($n - 1$). It is implemented as an array.¹⁴

¹⁴ insert figures

4 Strings

4.1 Information Theory

LECTURE 03/13

4.2 Data Structures

LECTURE 03/16

Items 1. - 3. are used for strings x_1, \dots, x_n where $x_i = (x_{i1} \ x_{i2} \ \dots)$ s.t. $x_{ij} \in A$ (alphabet); and items 4. - 6. are used for text $T = (t_1 \ \dots \ t_n)$ where $t_i \in A$.

1. TRIE: a tree-based data structure for storing strings, where each string is a path in the k -ary tree (each edge is a letter from the k -alphabet). The leaf stores the string's index.
2. PATRICIA TRIE: a compact trie, where all the nodes with one child are collapsed.
3. DIGITAL SEARCH TREE (DST): a k -ary tree where one inserts x_1, x_2, \dots, x_n in turn, and each string occupies the first available slot in a method like a binary search tree.
4. SUFFIX TRIE: a trie containing the n suffixes of the given text T .
5. SUFFIX TREE: a PATRICIA version of the suffix trie.
6. SUFFIX ARRAY: a sorted array containing a lexicographical ordering of the suffixes, i.e. with leaf numbers in the suffix trie seen in preorder. One can then locate strings by binary search $O(\log n)$.

Data structures used for a trie:

- (a) array of children (time efficient)
- (b) linked list of children (space efficient)

Show that for a k -ary PATRICIA, the number of nodes is less than $2n - 1$

$$\begin{aligned} T &= T[1] \ \dots \ T[n] = x_1 \\ T[2] \ \dots \ T[n] &= x_2 \\ &\dots \\ T[n] &= x_n \end{aligned}$$

4.3 String Matching

Algorithm by Knuth, Morris and Pratt given a text T of length n and pattern P of length m where $m \ll n$ which runs in time $O(m + n) = O(n)$ as opposed to the naive $O(mn)$.

KMP ALGORITHM(T, P)

- 1 preprocess P in time $O(m)$
- 2 actual string matching with T in time $O(n)$